

ANALYZING US BIRTH DATA FROM 2016 TO 2021

Eyram Kueviakoe

Professor Nasheb Ismaily

May 30, 2024

Bellevue University
DSC 650 – Big Data

Outline

Introduction	- 2-
Data Source	- 2-
Presentation of the dataset	- 2-
Data Ingestion	- 3-
Loading data into HDFS	- 3-
Using Hive	- 5-
Using PySpark	- 9-
Conclusion	- 12-

Introduction

Birth rates are a key indicator of population growth and can impact various sectors like healthcare and education. Understanding birth trends is important to improve the educational system and healthcare. It is also useful in making policies.

This project analyzes birth data from 2016 to 2021 in the US to discover trends and patterns. We will be focusing on how the educational level of parents affects birth rates. During this project, we will use big data tools to handle and analyze the dataset efficiently. We will use HDFS to store the data, spark to process the data, and Hive to query the data.

Data source

The dataset used for this project was obtained from Kaggle. It contains detailed information about the birth records from 2016 to 2021. It includes information on the number of births grouped by year, state and the educational level of the parents. The goal of the project is to use big data tools to analyze and query the data.

The link to the dataset is: https://www.kaggle.com/datasets/danbraswell/temporary-us-births?select=us_births_2016_2021.csv

Presentation of the dataset

Columns description:

State: Full name of the state

State Abbreviation: 2-character abbreviation of the state.

Year: The 4-digit year in which the births were recorded.

Gender: The gender of the baby (Male/Female).

Education Level of Mother: The education level of the mother.

Education Level Code: The corresponding code for the education level of the mother

Number of Births: The number of births for the specified category.

Average Age of Mother (years): The average age of the mother in the specified category.

Average Birth Weight (g): The average birth weight in grams for the specified category.

Showing a sample of the dataset (with relevant columns)

State	Year	Gender	Education Level	Number of Births
Vermont	2019	M	Master's degree (MA, MS, MEng, MEd, MSW, MBA)	403
Pennsylvania	2018	M	Doctorate (PhD, EdD) or Professional Degree (MD, DDS, DVM, LLB, JD)	2312
New Jersey	2020	F	9th through 12th grade with no diploma	2688
Texas	2021	M	Unknown or Not Stated	864
Oklahoma	2018	M	Some college credit, but not a degree	5875
Utah	2018	F	Some college credit, but not a degree	5579
New Mexico	2021	F	Associate degree (AA, AS)	1015
Washington	2020	M	Master's degree (MA, MS, MEng, MEd, MSW, MBA)	4652
Indiana	2017	M	9th through 12th grade with no diploma	4446
Vermont	2017	F	High school graduate or GED completed	717

Data Ingestion

The first step of this project is to import the dataset into HDFS (Hadoop Distributed File System). To achieve this, we will connect to the HDFS instance in our google cloud virtual machine and upload the csv file containing the data.

Step 1: We download the csv file from Kaggle and upload it to our github repository.

Step 2: Uploading the csv file into our virtual machine using the command

`wget https://raw.githubusercontent.com/kueyram/dsc650/main/us_births_2016_2021.csv`

```
rsa-key-20240315@dsc650-kueviakoe: ~
rsa-key-20240315@dsc650-kueviakoe:~$ wget https://github.com/kueyram/dsc650/blob/main/us_births_2016_2021.csv
--2024-06-01 18:39:14-- https://github.com/kueyram/dsc650/blob/main/us_births_2016_2021.csv
Resolving github.com (github.com)... 140.82.114.4
Connecting to github.com (github.com)|140.82.114.4|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: unspecified [text/html]
Saving to: 'us_births_2016_2021.csv'

us_births_2016_2021.csv                               [  <=>]

2024-06-01 18:39:15 (2.10 MB/s) - 'us_births_2016_2021.csv' saved [801983]

rsa-key-20240315@dsc650-kueviakoe:~$ ls
dsc650-infra  us_births_2016_2021.csv  world_population_data.csv
rsa-key-20240315@dsc650-kueviakoe:~$
```

Fig1: Downloading the csv file onto the virtual machine

Step 3: Starting the Docker container and then accessing the master container

```
rsa-key-20240315@dsc650-kueviakoe: ~
rsa-key-20240315@dsc650-kueviakoe:~$ wget https://raw.githubusercontent.com/kueyram/dsc650/main/us_births_2016_2021.csv
--2024-06-01 20:16:38-- https://raw.githubusercontent.com/kueyram/dsc650/main/us_births_2016_2021.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.111.133, 185.199.110.133, 185.199.109.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.111.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 416875 (407K) [text/plain]
Saving to: 'us_births_2016_2021.csv'

us_births_2016_2021.csv                               100%[=====]

2024-06-01 20:16:38 (17.5 MB/s) - 'us_births_2016_2021.csv' saved [416875/416875]

rsa-key-20240315@dsc650-kueviakoe:~$ ls
dsc650-infra  us_births_2016_2021.csv  world_population_data.csv
rsa-key-20240315@dsc650-kueviakoe:~$
```

Fig2: Starting docker container

Step 4: Load the csv file into HDFS

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase
bash-5.0# hdfs dfs -put /data/us_births.2016.2021.csv
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/program/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/tez/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
2024-06-01 18:49:17,733 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

Fig3: Loading the csv file into hdfs

Step 5: Checking if the file was successfully uploaded.

```

rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase
bash-5.0# hdfs dfs -ls
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/program/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/tez/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
2024-06-01 18:53:14,990 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x  - root supergroup          0 2024-06-01 18:45 .hiveJars
-rw-r--r--  1 root supergroup    675514 2024-06-01 18:49 us_births_2016_2021.csv
bash-5.0#

```

Fig4: Showing the csv was uploaded into hdfs

Create table and load data in Hive

Step6: Let's start a hive session and create a table using this command

```

CREATE TABLE birth_data(
`State` STRING,
`State Abbreviation` STRING,
`Year` INT,
`Gender` STRING,
`Education Level of Mother` STRING,
`Education Level Code` INT,
`Number of Births` INT,
`Average Age` FLOAT,
`Average Birth Weight (g)` FLOAT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE
tblproperties("skip.header.line.count"="1");

```

```

rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase
bash-5.0# hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/program/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/tez/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 7430a4bc-dfcc-4f43-a06a-331adda7a164

Logging initialized using configuration in file:/usr/program/hive/conf/hive-log4j2.properties Async: true
Hive Session ID = d787d6ad-9543-494b-9ed5-7b9ce3736655
2024-06-01 19:53:29,154 INFO [Tez session start thread] client.RMProxy: Connecting to ResourceManager at master/172.28.1.1:8032
2024-06-01 19:53:29,808 INFO [pool-7-thread-1] client.RMProxy: Connecting to ResourceManager at master/172.28.1.1:8032
hive> CREATE TABLE birth_data(
  > 'State' STRING,
  > 'State Abbreviation' STRING,
  > 'Year' INT,
  > 'Gender' STRING,
  > 'Education Level of Mother' STRING,
  > 'Education Level Code' INT,
  > 'Number of Births' INT,
  > 'Average Age' FLOAT,
  > 'Average Birth Weight (g)' FLOAT)
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY ','
  > STORED AS TEXTFILE
  > tblproperties("skip.header.line.count"="1");
OK
Time taken: 2.206 seconds

```

Fig5: Accessing Hive and creating a table to store the data

Step 7: Loading the data into the Hive table:

```

rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase
bash-5.0# hadoop fs -put /data/us_births_2016_2021.csv /
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/program/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/tez/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
2024-06-01 19:59:20,305 WARN util.NativeCodeLoader: Unable to load native-hadoop library for optimization: java.io.IOException: /usr/lib/libhadoop.so.1: cannot open shared object file: No such file or directory
bash-5.0# hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/program/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/program/tez/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 3e271035-dc3e-46ea-a6c9-f0f1de6c8496

Logging initialized using configuration in file:/usr/program/hive/conf/hive-log4j2.properties Async: true
Hive Session ID = bf1d3e18-7bd1-4001-b4bb-248daa82005c
2024-06-01 19:59:38,462 INFO [Tez session start thread] client.RMProxy: Connecting to ResourceManager at master/172.28.1.1:8032
2024-06-01 19:59:39,036 INFO [pool-7-thread-1] client.RMProxy: Connecting to ResourceManager at master/172.28.1.1:8032
hive> LOAD DATA INPATH '/us_births_2016_2021.csv' INTO TABLE birth_data
Loading data to table default.birth_data
OK
Time taken: 2.433 seconds
hive>

```

Fig 6: Loading the data into the Hive table

Step 8: Running queries on the data

- Query 1: Let's count the number of rows in the dataset

SELECT COUNT() AS NumberRows FROM birth_data;*

```
hive> SELECT COUNT(*) AS NumberRows FROM birth_data;
2024-06-01 21:29:28,998 INFO [e48a545d-70b7-430f-b5f6-dc12458a9b1b main] reducesink.VectorReduce
52c27d53
Query ID = root_20240601212928_f5838230-df9d-42a7-b4d4-39f724b44ee9
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1717276800122_0002)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    1          1          0          0          0          0
Reducer 2 ..... container  SUCCEEDED    1          1          0          0          0          0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 6.89 s
-----
OK
5496
Time taken: 7.935 seconds, Fetched: 1 row(s)
hive> █
```

Fig 7: Number of rows in the dataset

We have 5496 rows in the dataset

- Query 2: Let's count the number of births in the dataset

SELECT SUM(`Number of Births`) AS NumberBirths FROM birth_data;

```
hive> SELECT SUM(`Number of Births`) AS NumberBirths FROM birth_data;
2024-06-01 21:36:12,645 INFO [e48a545d-70b7-430f-b5f6-dc12458a9b1b main] reducesink.VectorReduceSink
7c93e69
Query ID = root_20240601213612_4e7167ee-5fa0-4edb-8e76-6d8cadea23bf
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1717276800122_0002)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    1          1          0          0          0          0
Reducer 2 ..... container  SUCCEEDED    1          1          0          0          0          0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 5.09 s
-----
OK
8889084
Time taken: 6.0 seconds, Fetched: 1 row(s)
hive> █
```

Fig8: Number of births in the dataset

There are 8889084 births in the dataset

- Query 3: Number of births in each state in 2020

```
SELECT State, SUM(`Number of Births`) AS total_births_2020
FROM birth_data
WHERE Year = 2020
GROUP BY State;
```

```
Alabama 27014
Alaska 3953
Arizona 32771
Arkansas 16681
California 173017
Colorado 19475
Connecticut 9722
Delaware 4386
District of Columbia 2978
Florida 88981
Georgia 54801
Hawaii 5725
Idaho 7695
Illinois 45170
Indiana 35185
Iowa 13224
Kansas 12718
Kentucky 22921
Louisiana 27872
Maine 3772
Maryland 23170
Massachusetts 17864
Michigan 39186
Minnesota 17472
Mississippi 15281
Missouri 26373
Montana 3845
Nebraska 7689
Nevada 16721
New Hampshire 3168
New Jersey 33249
New Mexico 9446
New York 77823
North Carolina 44131
North Dakota 3110
Ohio 53985
Oklahoma 21934
Oregon 13950
Pennsylvania 49634
Rhode Island 3447
South Carolina 21605
South Dakota 4413
Tennessee 33122
Texas 162202
Utah 13566
Vermont 1629
Virginia 34447
Washington 27883
West Virginia 7958
Wisconsin 21558
Wyoming 2179
Time taken: 7.486 seconds, Fetched: 51 row(s)
hive> █
```

Fig8: Number of births in each state in 2020

Query 4: Top 5 education level with the highest births

```
SELECT `Education Level of Mother`, SUM(`Number of Births`) AS Number_Births
FROM birth_data
GROUP BY `Education Level of Mother`
ORDER BY Number_Births DESC
LIMIT 5;
```

```
hive> SELECT `Education Level of Mother`, SUM(`Number of Births`) AS Number_Births
> FROM birth_data
> GROUP BY `Education Level of Mother`
> ORDER BY Number_Births DESC
> LIMIT 5;
2024-06-02 00:12:14,996 INFO [2c39b775-8956-4e74-b14e-61b23a1169cd main] reducesink.VectorReduceSinkOb
nfo@laaaabdl
Query ID = root_20240602001213_2339664b-47b9-48d6-9a6f-3896df313233
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1717286975291_0001)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    1          1          0          0          0          0
Reducer 2 ..... container  SUCCEEDED    1          1          0          0          0          0
Reducer 3 ..... container  SUCCEEDED    1          1          0          0          0          0
-----
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 6.53 s
-----
OK
High school graduate or GED completed  5775918
9th through 12th grade with no diploma 2086382
8th grade or less 708850
Unknown or Not Stated 312426
"Associate degree (AA 3060
Time taken: 10.849 seconds, Fetched: 5 row(s)
hive>
```

Fig 9: Top 5 education level with the highest births

Using Spark

Step 8: Loading data into spark and creating a dataframe

```
birth_data_df = spark.read.format('csv').option("header","true").load('us_births_2016_2021.csv')
```

Let's verify the dataframe was created.

```
Birth_data_df.show()
```

```
>>> birth_data_df = spark.read.format('csv').option('header','true').load('us_births_2016_2021.csv')
>>> birth_data_df.show()
```

State	State Abbreviation	Year	Gender	Education Level of Mother	Education Level Code	Number of Births	Average Age of Mother (years)	Average Birth Weight (g)
Alabama	AL	2016	F	8th grade or less	1	1052	27.8	3116.9
Alabama	AL	2016	F	9th through 12th ...	2	3436	24.1	3040.0
Alabama	AL	2016	F	High school gradu...	3	8777	25.4	3080.0
Alabama	AL	2016	F	Some college cred...	4	6453	26.7	3121.9
Alabama	AL	2016	F	Associate degree ...	5	2227	28.9	3174.3
Alabama	AL	2016	F	Bachelor's degree...	6	4453	30.3	3239.0
Alabama	AL	2016	F	Master's degree (...)	7	1910	32.0	3263.5
Alabama	AL	2016	F	Doctorate (PhD, E...	8	487	33.1	3196.7
Alabama	AL	2016	F	Unknown or Not St...	-9	65	27.7	3083.9
Alabama	AL	2016	M	8th grade or less	1	1188	27.6	3232.9
Alabama	AL	2016	M	9th through 12th ...	2	3657	23.9	3121.2
Alabama	AL	2016	M	High school gradu...	3	5284	25.2	3197.9
Alabama	AL	2016	M	Some college cred...	4	6816	26.7	3252.1
Alabama	AL	2016	M	Associate degree ...	5	2460	29.0	3301.4
Alabama	AL	2016	M	Bachelor's degree...	6	4645	30.3	3376.1
Alabama	AL	2016	M	Master's degree (...)	7	1974	32.2	3358.2
Alabama	AL	2016	M	Doctorate (PhD, E...	8	511	32.8	3368.4
Alabama	AL	2016	M	Unknown or Not St...	-9	56	27.2	3107.7
Alabama	AL	2017	F	8th grade or less	1	1012	27.6	3139.6
Alabama	AL	2017	F	9th through 12th ...	2	3283	24.4	3040.6

only showing top 20 rows

Fig 9: Checking that the dataframe was created and populated

Step 9: Let's remove the redundant columns from the dataframe

The dataframe has a column called State and another column which has the state names abbreviated. It also has the educational level and the code that corresponds to the educational level.

We will drop State Abbreviation, and Education Level Code

```
birth_data_cleaned = birth_data_df.drop("State Abbreviation", "Education Level Code")
```

Checking the new dataframe

```
Birth_data_cleaned.show()
```

```
>>> birth_data_cleaned.show()
+-----+-----+-----+-----+-----+-----+-----+
| State|Year|Gender|Education Level of Mother|Number of Births|Average Age of Mother (years)|Average Birth Weight (g)|
+-----+-----+-----+-----+-----+-----+-----+
|Alabama|2016|F|8th grade or less|1052|27.8|3116.9|
|Alabama|2016|F|9th through 12th ...|3436|24.1|3040.0|
|Alabama|2016|F|High school gradu...|8777|25.4|3080.0|
|Alabama|2016|F|Some college cred...|6453|26.7|3121.9|
|Alabama|2016|F|Associate degree ...|2227|28.9|3174.3|
|Alabama|2016|F|Bachelor's degree...|4453|30.3|3239.0|
|Alabama|2016|F|Master's degree (...|1910|32.0|3263.5|
|Alabama|2016|F|Doctorate (PhD, E...|487|33.1|3196.7|
|Alabama|2016|F|Unknown or Not St...|65|27.7|3083.9|
|Alabama|2016|M|8th grade or less|1188|27.6|3232.9|
|Alabama|2016|M|9th through 12th ...|3657|23.9|3121.2|
|Alabama|2016|M|High school gradu...|9284|25.2|3197.9|
|Alabama|2016|M|Some college cred...|6516|26.7|3252.1|
|Alabama|2016|M|Associate degree ...|2460|29.0|3301.4|
|Alabama|2016|M|Bachelor's degree...|4645|30.3|3376.1|
|Alabama|2016|M|Master's degree (...|1974|32.2|3358.2|
|Alabama|2016|M|Doctorate (PhD, E...|511|32.8|3368.4|
|Alabama|2016|M|Unknown or Not St...|56|27.2|3107.7|
|Alabama|2017|F|8th grade or less|1012|27.6|3139.6|
|Alabama|2017|F|9th through 12th ...|3283|24.4|3040.6|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

Fig 10: New dataframe after redundant columns were removed

With the cleaned dataframe, we can create visualizations and find possible correlations

between the number of births and the education level of the mother.

Step 10: Number of births grouped by education level

```
births_by_education_level = birth_data_cleaned.groupBy('Education Level of Mother').sum('Number of Births')
```

```
births_by_education_level = births_by_education_level.withColumnRenamed('sum(Number of Births)', 'Number of Births')
```

```
births_by_education_level.show()
```

```
>>> births_by_education_level = birth_data_cleaned.groupBy('Education Level of Mother').sum('Number of Births')
births_by_education_level = births_by_education_level.withColumnRenamed('sum(Number of Births)', 'Number of Births')
births_by_education_level.show()>>> births_by_education_level = births_by_education_level.withColumnRenamed('sum(Number of Births)', 'Number of Births')
>>> births_by_education_level.show()
+-----+-----+
|Education Level of Mother|Number of Births|
+-----+-----+
|9th through 12th ...|2086382|
|8th grade or less|708850|
|Unknown or Not St...|312426|
|Master's degree (...|2161046|
|Associate degree ...|1867700|
|Some college cred...|4425269|
|High school gradu...|5775918|
|Bachelor's degree...|4653184|
|Doctorate (PhD, E...|627705|
+-----+-----+
```

Fig 11: Births grouped by educational level

Conclusion

After ingesting the data into HDFS, we were able to define our table schema in Hive and import the data into the table. This gave us the opportunity to use HiveQL to run different queries on the table.

For advanced data manipulation and transformation, we use PySpark to create a dataframe that can be used to create visualization and graphs. PySpark can be used for data analysis, complex transformation, and machine learning.

The analysis of the data shows that as the mother's education level rises, there is a decrease in the number of births. This means that there is a possible correlation between higher education among mothers and reduced fertility rates. This can be used by healthcare professionals and educators to promote family planning.

In the future, we could include real-time data processing to ensure the analysis is up-to-date and responsive to changing trends. We could also use machine learning to make predictions. Finally, we could implement advanced visualization and interactive dashboards to present and communicate our findings.