

DSC 650

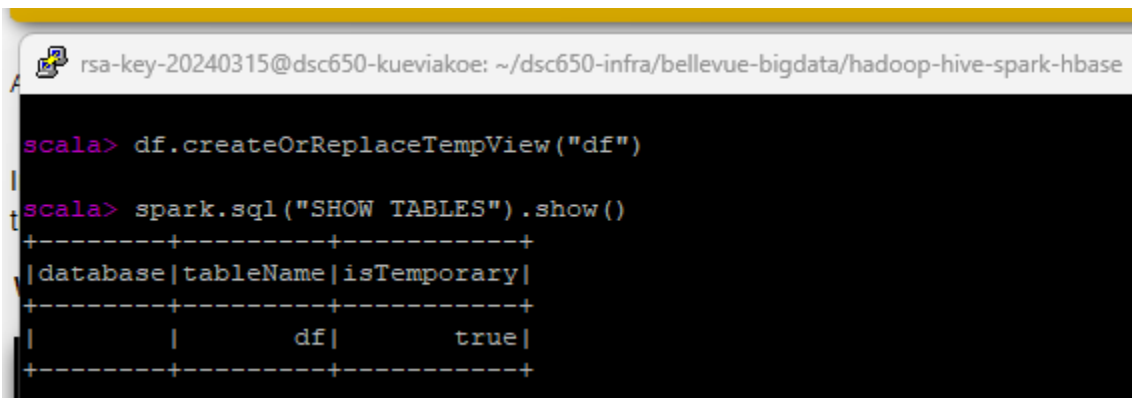
Week 5 Assignment

Eyram Kueviakoe

April 9, 2024

Screenshot of the results obtained from the SparkSQL commands in Scala

```
spark.sql("SHOW TABLES").show()
```

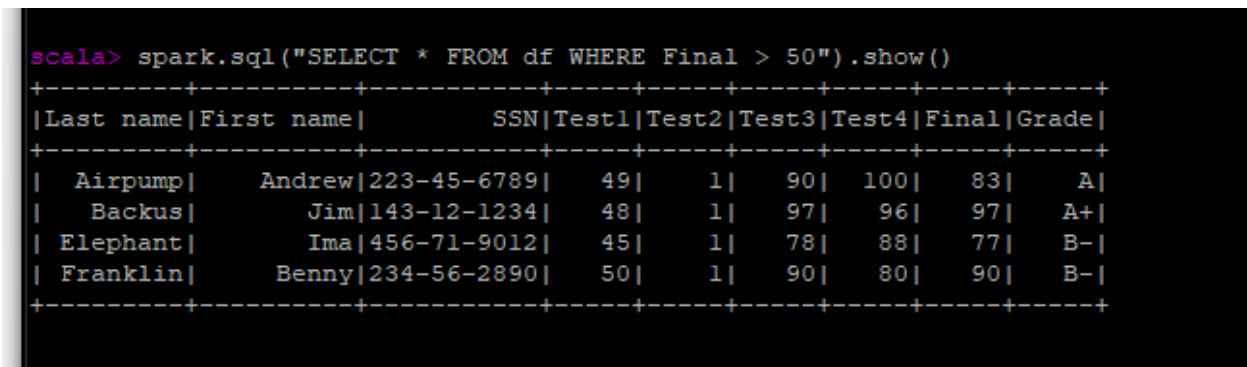


The screenshot shows a terminal window with a title bar indicating the user is 'rsa-key-20240315@dsc650-kueviakoe' and the current directory is '~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase'. The terminal displays the following commands and output:

```
scala> df.createOrReplaceTempView("df")
scala> spark.sql("SHOW TABLES").show()
```

database	tableName	isTemporary
	df	true

```
spark.sql("SELECT * FROM df WHERE Final > 50").show()
```



The screenshot shows the terminal window displaying the results of the query 'SELECT \* FROM df WHERE Final > 50'. The output is a table with 9 columns: Last name, First name, SSN, Test1, Test2, Test3, Test4, Final, and Grade. The results are as follows:

Last name	First name	SSN	Test1	Test2	Test3	Test4	Final	Grade
Airpump	Andrew	223-45-6789	49	1	90	100	83	A
Backus	Jim	143-12-1234	48	1	97	96	97	A+
Elephant	Ima	456-71-9012	45	1	78	88	77	B-
Franklin	Benny	234-56-2890	50	1	90	80	90	B-

```
spark.sql("SELECT * FROM grades").show()
```

scala> spark.sql("SELECT \* FROM df").show()

Last name	First name	SSN	Test1	Test2	Test3	Test4	Final	Grade
Alfalpa	Aloysius	123-45-6789	40	90	100	83	49	D-
Alfred	University	123-12-1234	41	97	96	97	48	D+
Gerty	Gramma	567-89-0123	41	80	60	40	44	C
Android	Electric	087-65-4321	42	23	36	45	47	B-
Bumpkin	Fred	456-78-9012	43	78	88	77	45	A-
Rubble	Betty	234-56-7890	44	90	80	90	46	C-
Noshow	Cecil	345-67-8901	45	11	-1	4	43	F
Buff	Bif	632-79-9939	46	20	30	40	50	B+
Airpump	Andrew	223-45-6789	49	1	90	100	83	A
Backus	Jim	143-12-1234	48	1	97	96	97	A+
Carnivore	Art	565-89-0123	44	1	80	60	40	D+
Dandy	Jim	087-75-4321	47	1	23	36	45	C+
Elephant	Ima	456-71-9012	45	1	78	88	77	B-
Franklin	Benny	234-56-2890	50	1	90	80	90	B-
George	Boy	345-67-3901	40	1	11	-1	4	B
Heffalump	Harvey	632-79-9439	30	1	20	30	40	C

scala>

scala>

Screenshot of your 3 other SQL query results

Query 1: Top 5 students with the highest scores in Test 1

```
spark.sql("SELECT `Last name`, `First name`, Test1 FROM df ORDER BY Test1 DESC LIMIT 5").show()
```

scala> spark.sql("SELECT `Last name`, `First name`, Test1 FROM df ORDER BY Test1 DESC LIMIT 5").show()

Last name	First name	Test1
Franklin	Benny	50
Airpump	Andrew	49
Backus	Jim	48
Dandy	Jim	47
Buff	Bif	46

scala> █

## Query 2: Students with highest Final exam score

`spark.sql("SELECT `Last name`, `First name`, MAX(Final) AS Highest_Final FROM df GROUP BY `Last name`, `First name` ORDER BY Highest_Final DESC").show()`

```
scala> spark.sql("SELECT `Last name`, `First name`, MAX(Final) AS Highest_Final FROM df GROUP BY `Last name`, `First name` ORDER BY Highest_Final DESC").show()
+-----+-----+-----+
|Last name|First name|Highest_Final|
+-----+-----+-----+
|Backus|Jim|97|
|Franklin|Benny|90|
|Airpump|Andrew|83|
|Elephant|Ima|77|
|Buff|Bif|50|
|Alfalfa|Aloysius|49|
|Alfred|University|48|
|Android|Electric|47|
|Rubble|Betty|46|
|Bumpkin|Fred|45|
|Dandy|Jim|45|
|Gerty|Gramma|44|
|Noshov|Cecil|43|
|Heffalump|Harvey|40|
|Carnivore|Art|40|
|George|Boy|4|
+-----+-----+-----+
scala>
```

## Query 3: List of students who scored less than the average final exam score

`spark.sql("SELECT `Last name`, `First name`, Final FROM df WHERE Final < (SELECT AVG(Final) FROM df)").show()`

```
scala> spark.sql("SELECT `Last name`, `First name`, Final FROM df WHERE Final < (SELECT AVG(Final) FROM df)").show()
+-----+-----+-----+
|Last name|First name|Final|
+-----+-----+-----+
|Alfalfa|Aloysius|49|
|Alfred|University|48|
|Gerty|Gramma|44|
|Android|Electric|47|
|Bumpkin|Fred|45|
|Rubble|Betty|46|
|Noshov|Cecil|43|
|Buff|Bif|50|
|Carnivore|Art|40|
|Dandy|Jim|45|
|George|Boy|4|
|Heffalump|Harvey|40|
+-----+-----+-----+
scala>
```

Screenshot of the results obtained from the SparkSQL commands in Python.

```
spark.sql('SHOW TABLES').show()
```

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hiv
>>> df.createOrReplaceTempView('df')
>>> spark.sql('SHOW TABLES').show()
140045 [Thread-4] WARN  org.apache.hadoop.hive.conf.HiveCon
140046 [Thread-4] WARN  org.apache.hadoop.hive.conf.HiveCon
140077 [Thread-4] WARN  org.apache.spark.sql.hive.client.Hi
+-----+-----+-----+
|database|tableName|isTemporary|
+-----+-----+-----+
|         |         df|         true|
+-----+-----+-----+
```

```
spark.sql('SELECT * FROM df WHERE Final > 50').show()
```

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase
+-----+-----+-----+
|         |         df|         true|
+-----+-----+-----+
>>> spark.sql('SELECT * FROM df WHERE Final > 50').show()
+-----+-----+-----+-----+-----+-----+-----+-----+
|Last name|First name|      SSN|Test1|Test2|Test3|Test4|Final|Grade|
+-----+-----+-----+-----+-----+-----+-----+-----+
|  Airpump|   Andrew|223-45-6789|  49|   1|  90| 100|  83|   A|
|   Backus|    Jim|143-12-1234|  48|   1|  97|  96|  97|  A+|
| Elephant|    Ima|456-71-9012|  45|   1|  78|  88|  77|  B-|
| Franklin|   Benny|234-56-2890|  50|   1|  90|  80|  90|  B-|
+-----+-----+-----+-----+-----+-----+-----+-----+
```

`spark.sql('SELECT * FROM df').show()`

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase

>>> spark.sql('SELECT * FROM df').show()
+-----+-----+-----+-----+-----+-----+-----+-----+
|Last name|First name|SSN|Test1|Test2|Test3|Test4|Final|Grade|
+-----+-----+-----+-----+-----+-----+-----+-----+
|Alfalfa|Aloysius|123-45-6789|40|90|100|83|49|D-|
|Alfred|University|123-12-1234|41|97|96|97|48|D+|
|Gerty|Gramma|567-89-0123|41|80|60|40|44|C|
|Android|Electric|087-65-4321|42|23|36|45|47|B-|
|Bumpkin|Fred|456-78-9012|43|78|88|77|45|A-|
|Rubble|Betty|234-56-7890|44|90|80|90|46|C-|
|Noshov|Cecil|345-67-8901|45|11|-1|4|43|F|
|Buff|Bif|632-79-9939|46|20|30|40|50|B+|
|Airpump|Andrew|223-45-6789|49|1|90|100|83|A|
|Backus|Jim|143-12-1234|48|1|97|96|97|A+|
|Carnivore|Art|565-89-0123|44|1|80|60|40|D+|
|Dandy|Jim|087-75-4321|47|1|23|36|45|C+|
|Elephant|Ima|456-71-9012|45|1|78|88|77|B-|
|Franklin|Benny|234-56-2890|50|1|90|80|90|B-|
|George|Boy|345-67-3901|40|1|11|-1|4|B|
|Heffalump|Harvey|632-79-9439|30|1|20|30|40|C|
+-----+-----+-----+-----+-----+-----+-----+-----+
```

### Run 3 other SQL queries in the PySpark Shell

#### Query 1: Top 5 students with the highest scores in Test 1

`spark.sql("SELECT `Last name`, `First name`, Test1 FROM df ORDER BY Test1 DESC LIMIT 5").show()`

```
>>> spark.sql("SELECT `Last name`, `First name`, Test1 FROM df ORDER BY Test1 DESC LIMIT 5").show()
+-----+-----+-----+
|Last name|First name|Test1|
+-----+-----+-----+
|Franklin|Benny|50|
|Airpump|Andrew|49|
|Backus|Jim|48|
|Dandy|Jim|47|
|Buff|Bif|46|
+-----+-----+-----+

>>> █
```

### Query 2: Students with highest Final exam score

`spark.sql("SELECT `Last name`, `First name`, MAX(Final) AS Highest_Final FROM df GROUP BY `Last name`, `First name` ORDER BY Highest_Final DESC").show()`

```
>>> spark.sql("SELECT `Last name`, `First name`, MAX(Final) AS Highest_Final FROM df GROUP BY `Last name`, `First name` ORDER BY Highest_Final DESC").show()
+-----+-----+-----+
|Last name|First name|Highest_Final|
+-----+-----+-----+
|Backus|Jim|97|
|Franklin|Benny|90|
|Airpump|Andrew|83|
|Elephant|Ima|77|
|Buff|Bif|50|
|Alfalfa|Aloysius|49|
|Alfred|University|48|
|Android|Electric|47|
|Rubble|Betty|46|
|Bumpkin|Fred|45|
|Dandy|Jim|45|
|Gerty|Gramma|44|
|Noshov|Cecil|43|
|Heffalump|Harvey|40|
|Carnivore|Art|40|
|George|Boy|4|
+-----+-----+-----+
```

### Query 3: List of students who scored less than the average final exam score

`spark.sql("SELECT `Last name`, `First name`, Final FROM df WHERE Final < (SELECT AVG(Final) FROM df)").show()`

```
>>> spark.sql("SELECT `Last name`, `First name`, Final FROM df WHERE Final < (SELECT AVG(Final) FROM df)").show()
+-----+-----+-----+
|Last name|First name|Final|
+-----+-----+-----+
|Alfalfa|Aloysius|49|
|Alfred|University|48|
|Gerty|Gramma|44|
|Android|Electric|47|
|Bumpkin|Fred|45|
|Rubble|Betty|46|
|Noshov|Cecil|43|
|Buff|Bif|50|
|Carnivore|Art|40|
|Dandy|Jim|45|
|George|Boy|4|
|Heffalump|Harvey|40|
+-----+-----+-----+
>>> █
```

## 3- SparkSQL with custom data set

Our dataset from assignment 3 is world\_pop\_data.csv.

Loading data into Spark

```
val df = spark.read.format("csv").option("header", "true").load("/data/world_pop_data.csv")
df.createOrReplaceTempView("df")
```

### Query 1: 10 most populated countries in 2023

```
spark.sql("SELECT Country, Continent, Population_2023 FROM df ORDER BY Population_2023 DESC LIMIT 10").show()
```

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase

scala> spark.sql("SELECT Country, Continent, Population_2023 FROM df ORDER BY Population_2023 DESC LIMIT 10").show()
+-----+-----+-----+
| Country|Continent|Population_2023|
+-----+-----+-----+
| Vietnam|Asia|98858950|
| United States Vir...|North America|98750|
| Reunion|Africa|981796|
| United Arab Emirates|Asia|9516871|
| Belarus|Europe|9498238|
| Antigua and Barbuda|North America|94298|
| Fiji|Oceania|936375|
| Israel|Asia|9174520|
| Togo|Africa|9053799|
| Austria|Europe|8958960|
+-----+-----+-----+
```

### Query 2: What are the top 5 countries with highest density in 2023

```
spark.sql("SELECT Country, Continent, Population_2023/Area_km2 as Density FROM df ORDER BY Density DESC LIMIT 5").show()
```

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase

scala> spark.sql("SELECT Country, Continent, Population_2023/Area_km2 as Density FROM df ORDER BY Density DESC LIMIT 5").show()
+-----+-----+-----+
| Country|Continent|Density|
+-----+-----+-----+
| Macau|Asia|21402.70516717325|
| Monaco|Europe|17968.811881188118|
| Singapore|Asia|8471.440845070423|
| Hong Kong|Asia|6785.877717391304|
| Gibraltar|Europe|4807.058823529412|
+-----+-----+-----+
```

### Query 3: Find countries with a population greater than 100 million in 1970:

```
spark.sql("SELECT Country, Population_1970 FROM df WHERE Population_1970> 100000000").show()
```

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase

scala> spark.sql("SELECT Country, Population_1970 FROM df WHERE Population_1970> 100000000").show()
+-----+-----+
| Country|Population_1970|
+-----+-----+
| India|557501301|
| China|822534450|
| United States|200328340|
| Indonesia|115228394|
| Russia|130093010|
| Japan|105416839|
+-----+-----+
```

#### Query 4: Find countries with a population greater than 100 million in 2023:

```
spark.sql("SELECT Country, Population_2023 FROM df WHERE Population_2023 > 100000000").show()
```

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase

scala> spark.sql("SELECT Country, Population_2023 FROM df WHERE Population_2023 > 100000000").show()
+-----+-----+
|      Country|Population_2023|
+-----+-----+
|      India|      1428627663|
|      China|      1425671352|
|United States|      339996563|
|  Indonesia|      277534122|
|   Pakistan|      240485658|
|    Nigeria|      223804632|
|    Brazil|      216422446|
|Bangladesh|      172954319|
|    Russia|      144444359|
|    Mexico|      128455567|
|   Ethiopia|      126527060|
|    Japan|      123294513|
|Philippines|      117337368|
|    Egypt|      112716598|
|   DR Congo|      102262808|
+-----+-----+
```

#### Running the same queries using PySpark

##### Query 1: 10 most populated countries in 2023

```
spark.sql("SELECT Country, Continent, Population_2023 FROM df ORDER BY Population_2023 DESC LIMIT 10").show()
```

```
rsa-key-20240315@dsc650-kueviakoe: ~/dsc650-infra/bellevue-bigdata/hadoop-hive-spark-hbase

>>> df.createOrReplaceTempView('df')
>>> spark.sql("SELECT Country, Continent, Population_2023 FROM df ORDER BY Population_2023 DESC LIMIT 10").show()
+-----+-----+-----+
|      Country|Continent|Population_2023|
+-----+-----+-----+
|    Vietnam|      Asia|      98858950|
|United States Vir...|North America|      98750|
|    Reunion|      Africa|      981796|
|United Arab Emirates|      Asia|      9516871|
|    Belarus|      Europe|      9498238|
|Antigua and Barbuda|North America|      94298|
|    Fiji|      Oceania|      936375|
|    Israel|      Asia|      9174520|
|    Togo|      Africa|      9053799|
|    Austria|      Europe|      8958960|
+-----+-----+-----+
```



### Query 2: What are the top 5 countries with highest density in 2023

```
spark.sql("SELECT Country, Continent, Population_2023/Area_km2 as Density FROM df ORDER BY Density DESC LIMIT 5").show()
```

```
>>> spark.sql("SELECT Country, Continent, Population_2023/Area_km2 as Density FROM df ORDER BY Density DESC LIMIT 5").show()
+-----+-----+-----+
| Country|Continent|          Density|
+-----+-----+-----+
| Macau|      Asia| 21402.70516717325|
| Monaco|    Europe|17968.811881188118|
| Singapore|    Asia| 8471.440845070423|
| Hong Kong|    Asia| 6785.877717391304|
| Gibraltar| Europe| 4807.058823529412|
+-----+-----+-----+
```

### Query 3: Find countries with a population greater than 100 million in 1970:

```
spark.sql("SELECT Country, Population_1970 FROM df WHERE Population_1970> 100000000").show()
```

```
>>> spark.sql("SELECT Country, Population_1970 FROM df WHERE Population_1970> 100000000").show()
+-----+-----+
| Country|Population_1970|
+-----+-----+
| India|      557501301|
| China|      822534450|
| United States|    200328340|
| Indonesia|    115228394|
| Russia|    130093010|
| Japan|     105416839|
+-----+-----+
```

### Query 4: Find countries with a population greater than 100 million in 2023:

```
spark.sql("SELECT Country, Population_2023 FROM df WHERE Population_2023 > 100000000").show()
```

```
>>> spark.sql("SELECT Country, Population_2023 FROM df WHERE Population_2023 > 100000000").show()
+-----+-----+
| Country|Population_2023|
+-----+-----+
| India|    1428627663|
| China|    1425671352|
| United States|    339996563|
| Indonesia|    277534122|
| Pakistan|    240485658|
| Nigeria|    223804632|
| Brazil|    216422446|
| Bangladesh|    172954319|
| Russia|    144444359|
| Mexico|    128455567|
| Ethiopia|    126527060|
| Japan|    123294513|
| Philippines|    117337368|
| Egypt|    112716598|
| DR Congo|    102262808|
+-----+-----+
```