Week 6 Assignment: Hands-on with HBase

Objective:

Familiarize with the core functionalities of HBase, understand table creation, and querying data with a focus on data generation.

1. Environment Initialization

- Start by navigating to the required directory and initiating the Docker containers:

```
cd bellevue-bigdata
cd hadoop-hive-spark-hbase
docker-compose up -d
```

If you're using Google Cloud, remember to set up port forwarding as outlined in the previous assignments.

- Access the master container:

```
docker-compose exec master bash
```

2. Introduction to HBase

- Enter the HBase interactive shell:

```
hbase shell
```

3. Table Creation and Management

Exercise 1: Create a table named 'students' with a column family 'details'.

```
create 'students', 'details'
```

Deliverable: Screenshot of the table creation command and its output.

Exercise 2: Verify that the table has been created.

```
list
```

Deliverable: Screenshot of the tables listed in HBase.

4. Data Manipulation in HBase

Exercise 3: Add data to the 'students' table. Let's assume each student has a unique ID, a first name, and a last name.

```
put 'students', '1', 'details:firstName', 'John'
put 'students', '1', 'details:lastName', 'Doe'
```

Deliverable: Screenshot of the commands used to add data and their outputs.

Exercise 4: Query the data from the 'students' table to retrieve the details of the student with ID '1'.

```
get 'students', '1'
```

Deliverable: Screenshot of the query and its output.

5. Advanced HBase Features: Composite Row Key

Exercise 5: Create a table named 'orders' to store data about customer orders. Assume each order is uniquely identified by a composite key formed by combining the customer ID and order date (in the format YYYYMMDD).

```
create 'orders', 'orderDetails'
```

Exercise 6: Add sample data to the 'orders' table using the composite key:

```
put 'orders', '101:20230806', 'orderDetails:item', 'Laptop'
put 'orders', '102:20230806', 'orderDetails:item', 'Smartphone'
```

Exercise 7: Query the 'orders' table to retrieve details of all orders placed by the customer with ID '101'.

```
scan 'orders', {STARTROW => '101:', ENDROW => '101:~'}
```

This command will scan rows starting from '101:' to before '101:~' (tilde '~' is the next ASCII character after colon ':').

Deliverable: Screenshot of the commands used to query the data with composite key and their outputs.

6. Data Generation for HBase

Exercise 8: Generate random data for the 'students' table.

```
(2..100).each do |i|
  first_name = "Student#{i}"
  last_name = "LastName#{i}"
  put 'students', "#{i}", 'details:firstName', first_name
  put 'students', "#{i}", 'details:lastName', last_name
end
```

Exercise 9: Scan the 'students' table to verify data insertion.

```
scan 'students'
```

Deliverable: Screenshot of the commands used for data generation and their outputs.

Exercise 10: HBase Data Manipulation

Tasks:

1. **Update First Names:**
   - For students with IDs from `2` to `50`, change the first name prefix from `Student` to `Scholar`. For instance, `Student3` should become `Scholar3`.
2. **Add a Middle Name:**
   - For students with IDs from `51` to `75`, add a middle name column under the `details` column family. The middle name should follow the pattern `MidName#{i}`.
3. **Modify Last Names:**
   - For students with IDs from `76` to `100`, append `_Modified` to the last name. So, `LastName76` should be updated to `LastName76_Modified`.
4. **Bulk Delete:**
   - Delete all the details for students with IDs from `90` to `100`.
5. **Data Retrieval:**
   - After all modifications, retrieve and display the details for students with IDs `40`, `60`, `80`, and `90` to verify changes.

Deliverable:

- Screenshot of the commands used to implement the tasks above.
- Screenshots of the resulting output from a scan.

Shutting Down

Ensure all Docker containers are turned off with `docker-compose down` for each directory.
If you're using google cloud, please shut down your virtual machine to preserve cloud costs.